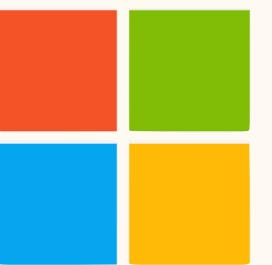


Aurora: A Foundation Model for the Earth System

Wessel Bruinsma

Microsoft Research AI for Science

**Building a Dutch AI-Earth System Modeling Community
KNMI, De Bilt, 29 Jan 2025**



The Aurora Team



Paris Perdikaris

University of Pennsylvania,
formerly MSR



Richard Turner

University of Cambridge,
formerly MSR



Max Welling

University of Amsterdam,
CuspAI, formerly MSR



Wessel Bruinsma

MSR



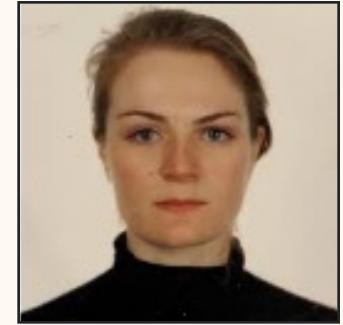
Anna Vaughan

University of Cambridge



Elizabeth Heider

Book tour, formerly MSR



Megan Stanley

MSR



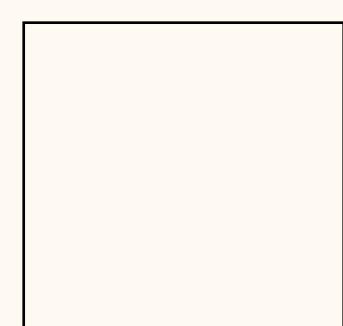
Johannes Brandstetter

JKU Linz, NXAI, formerly MSR



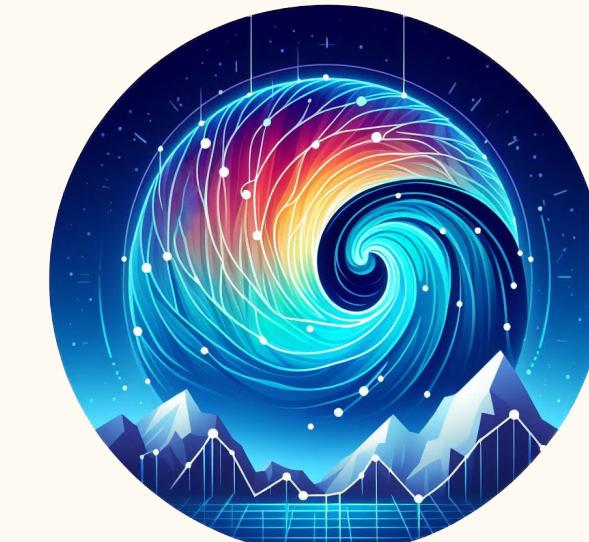
Ana Lučić

University of Amsterdam,
formerly MSR



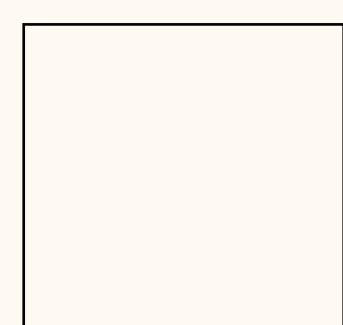
Patrick Garvan

Formerly MSR



Cristian Bodnar

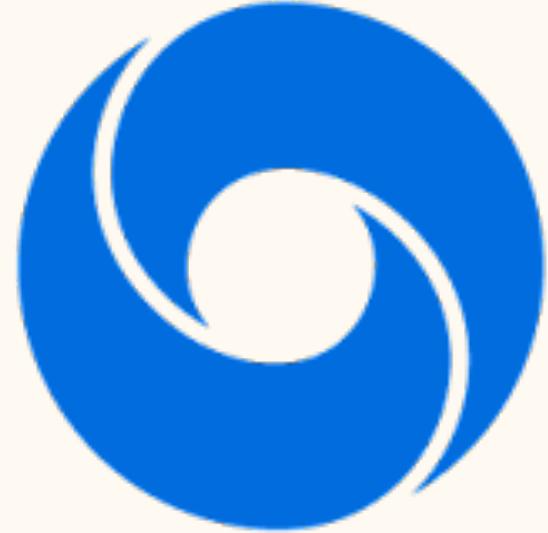
Silurian, formerly MSR



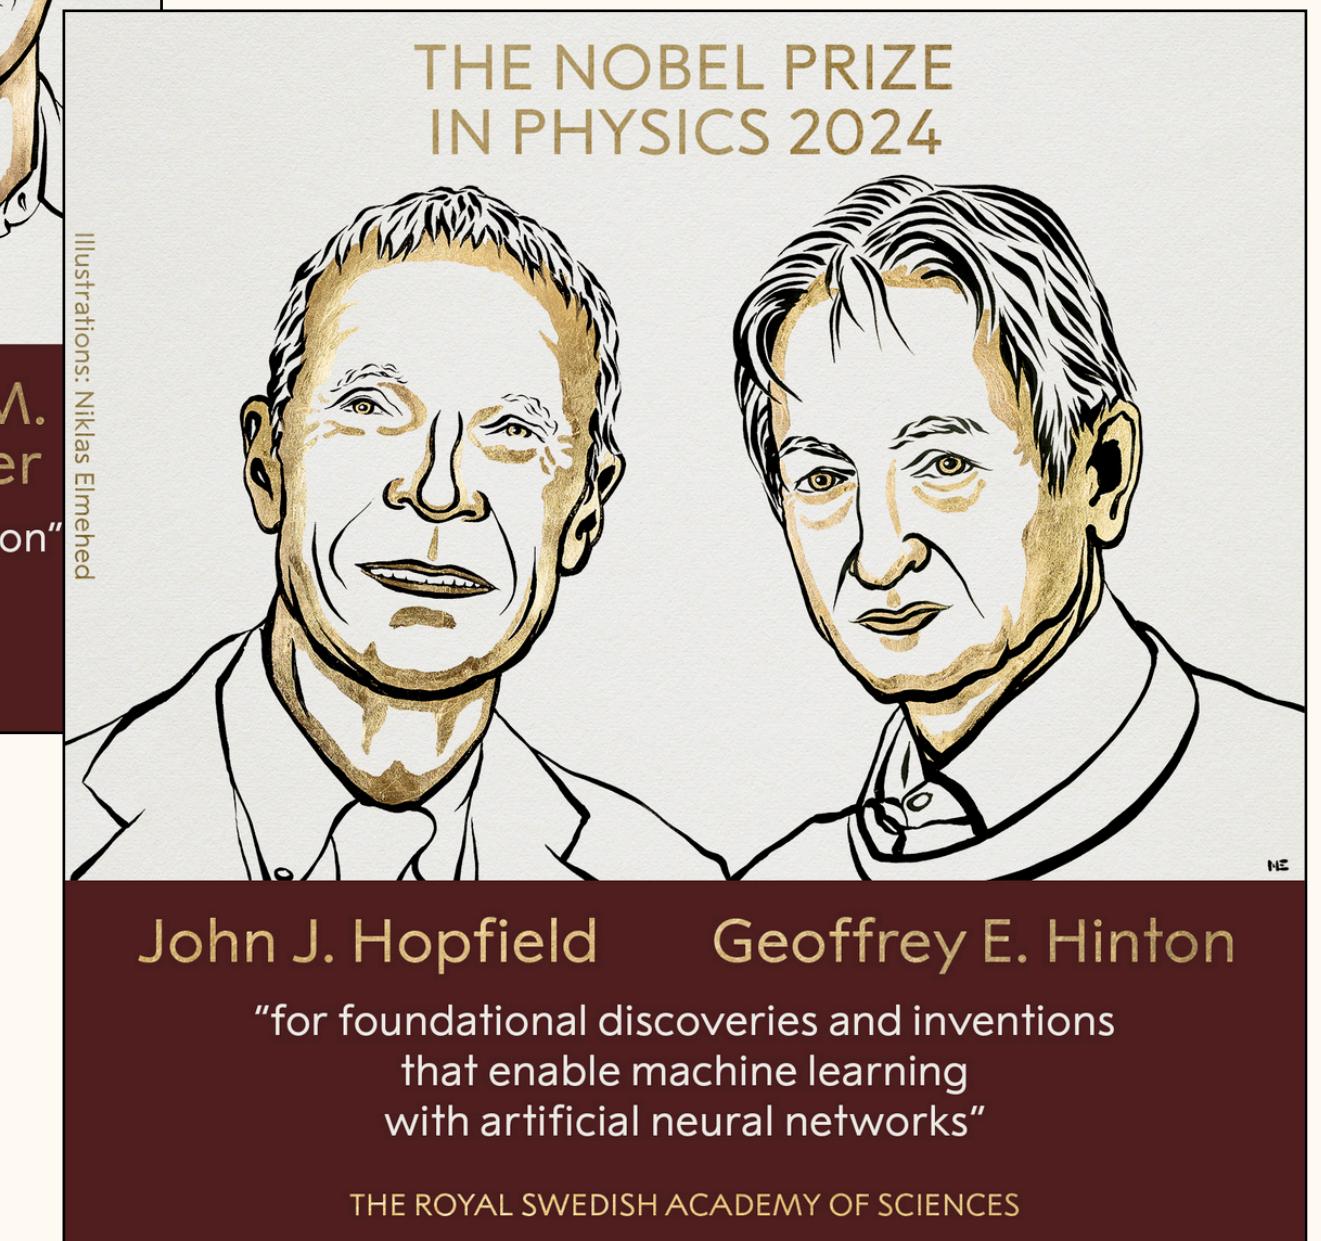
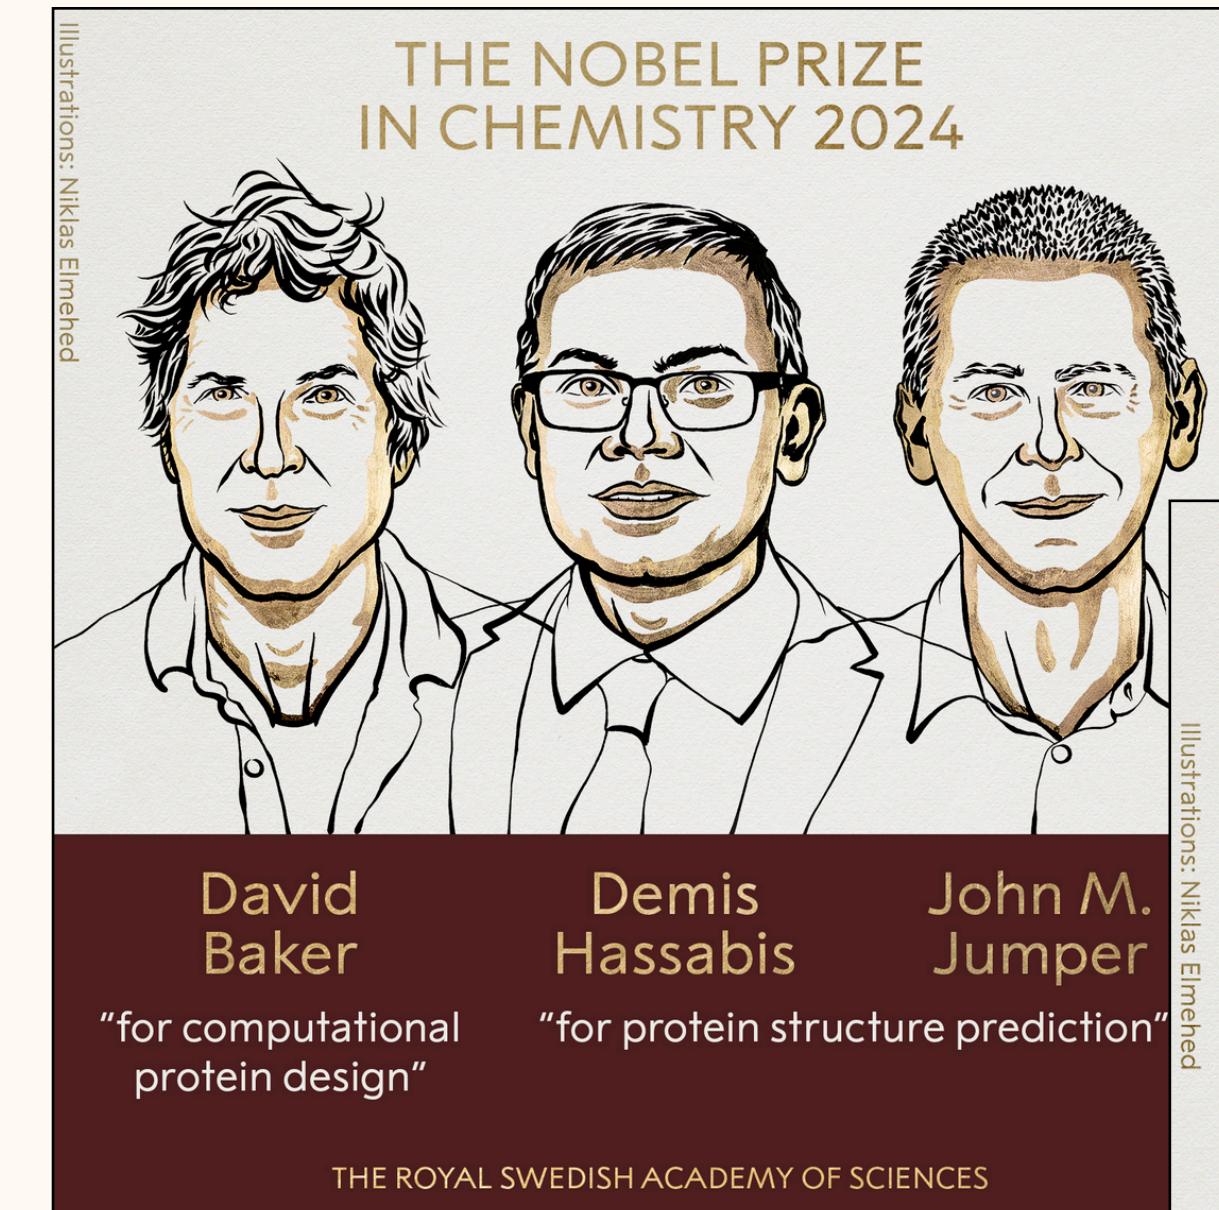
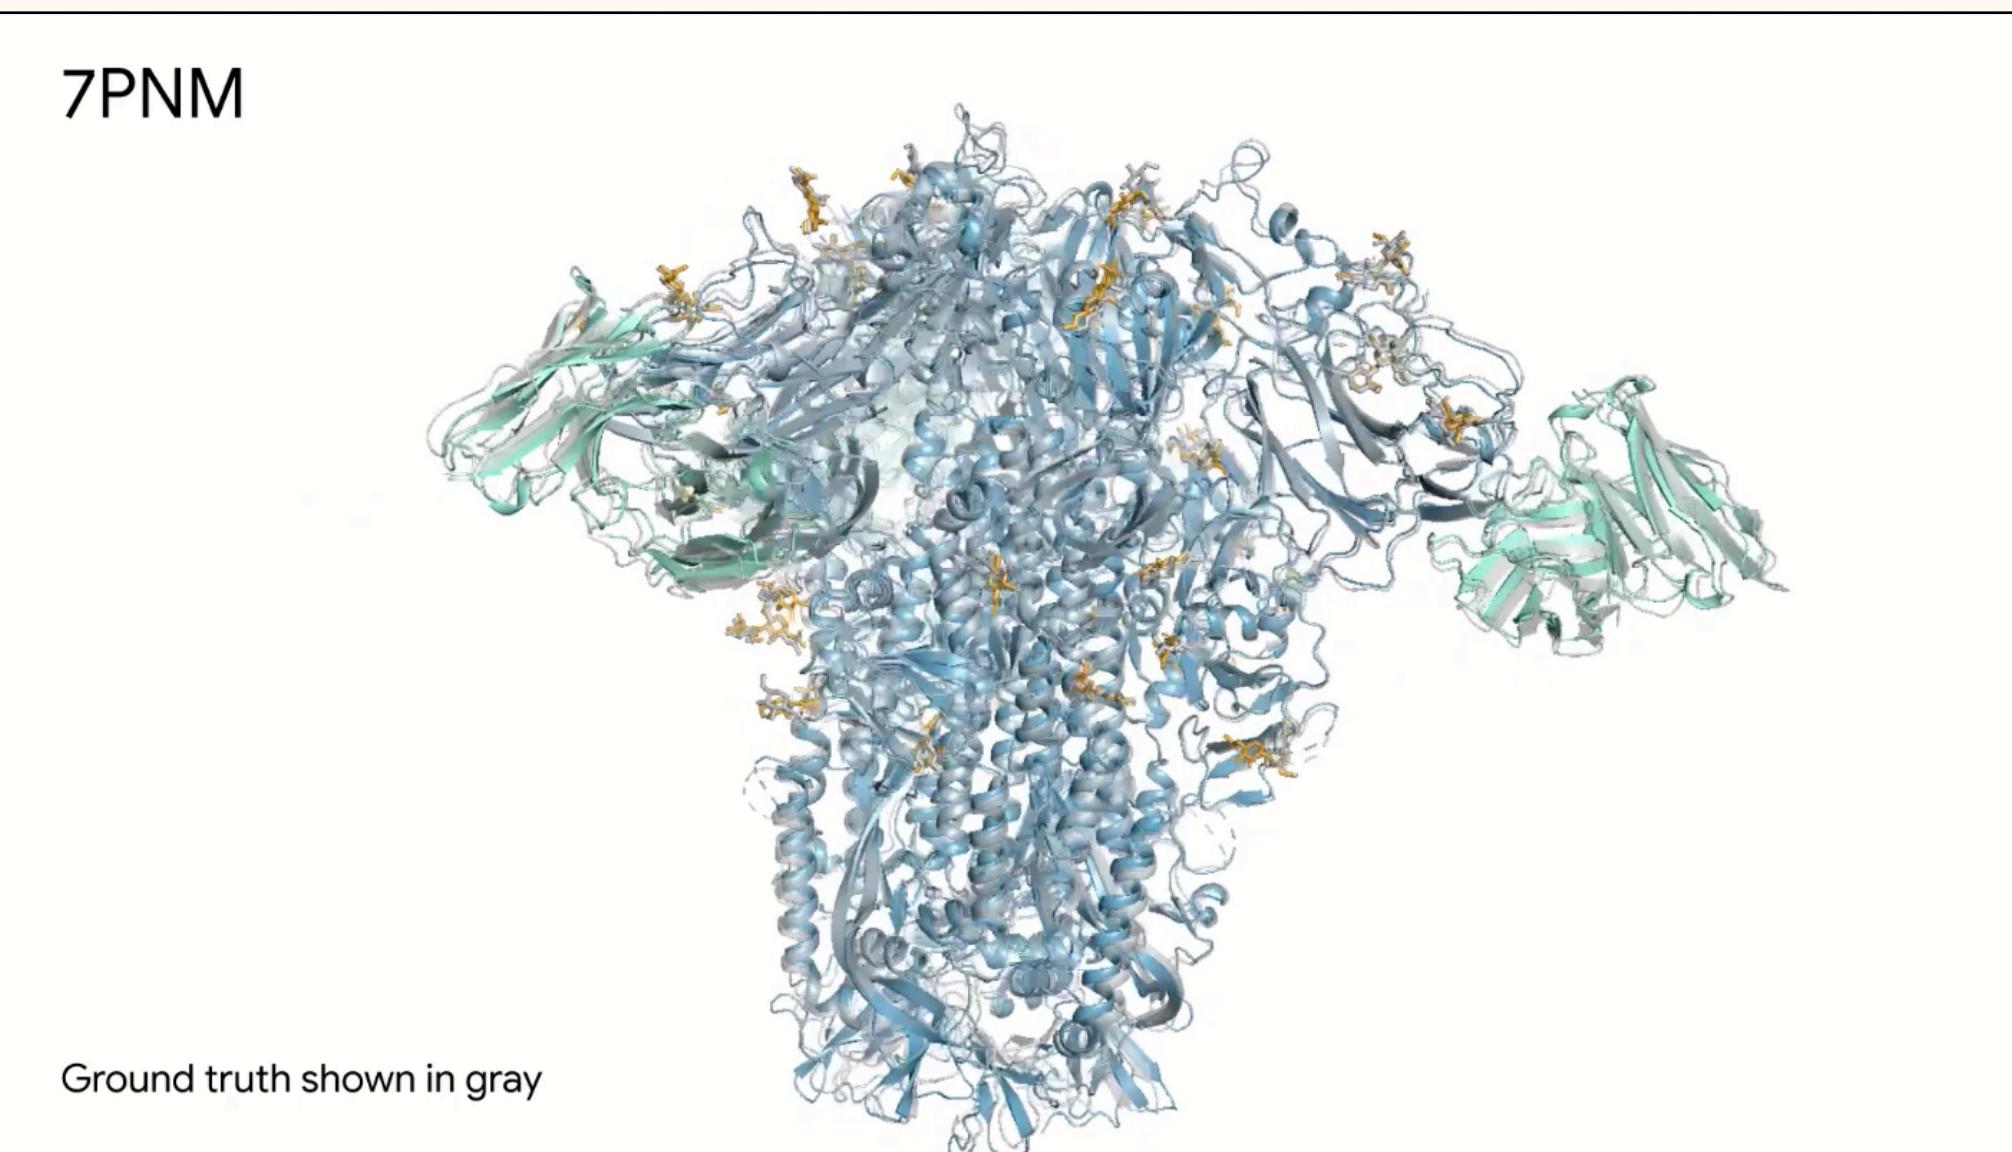
Maik Riechert

MSR

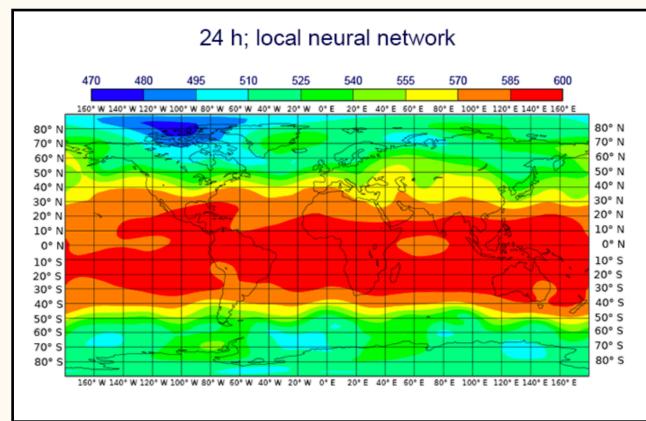
The AI Revolution in Science



AlphaFold
Protein folding



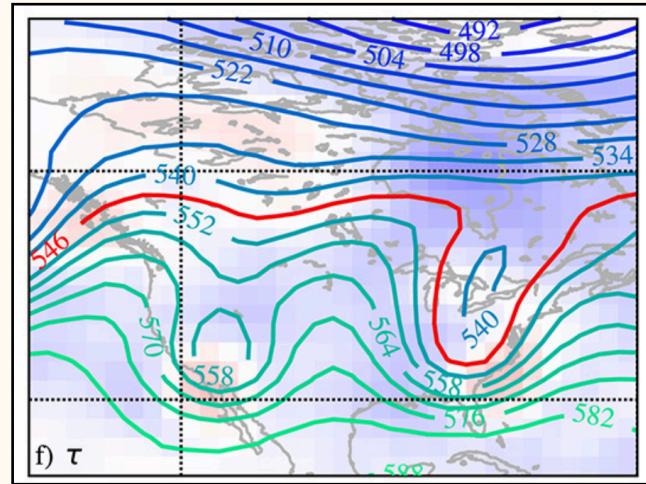
The AI Revolution in Weather Forecasting



2018

First serious efforts to compare AI models to physics baselines

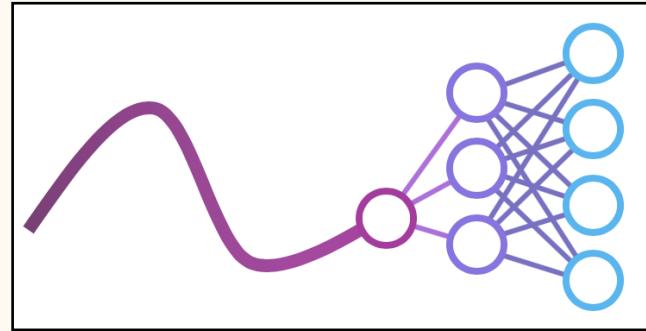
Dueben and Bauer (2018)



2019

AI models skillful to multiple days

Weyn et al. (2019)

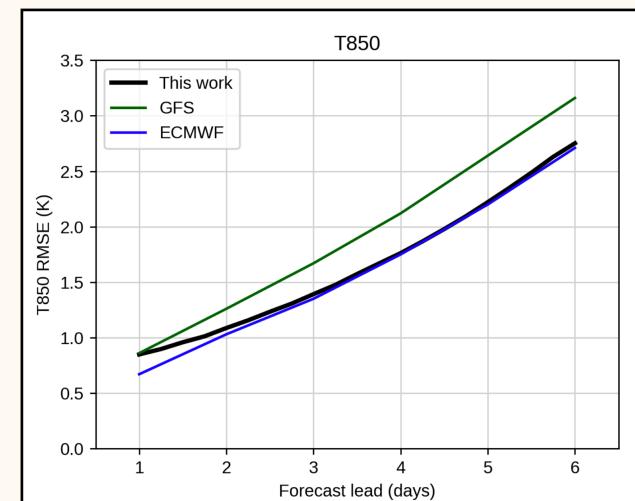


2020

WeatherBench starts to drive ML development

Rasp et al. (2018)

The AI Revolution in Weather Forecasting



2022

GNN outperforms GFS at 1°
Keisler (2022)



2022

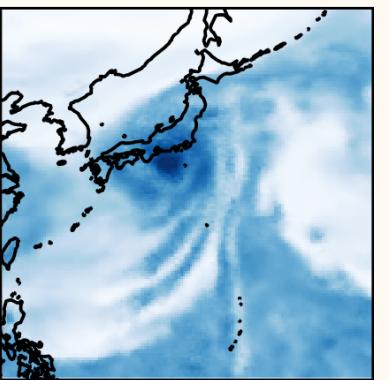
Pangu-Weather outperforms HRES at 0.25°
Bi et al. (2023)

The AI Revolution in Weather Forecasting



2022–2023

Tech companies start to work in this space



2023

GenCast outperforms IFS ensemble

Price et al. (2024)



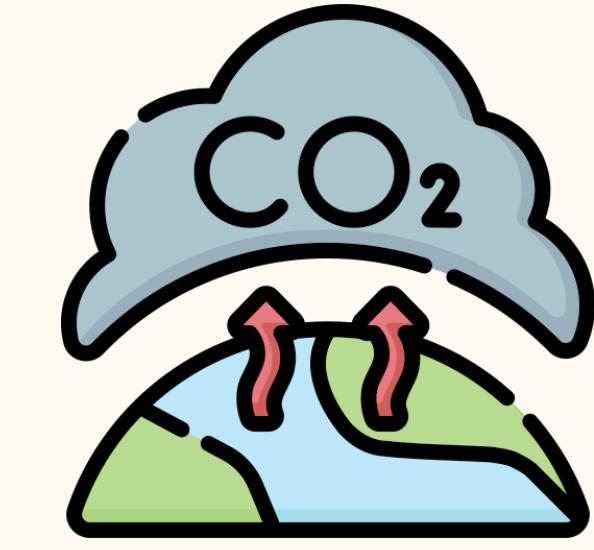
2024

ECMWF launches AIFS

What About Other Forecasting Tasks?



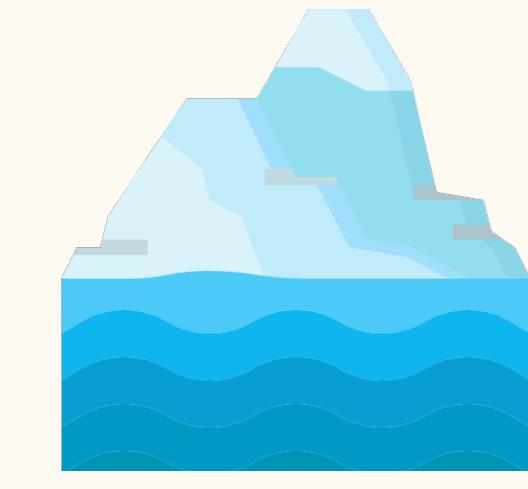
Air
pollution



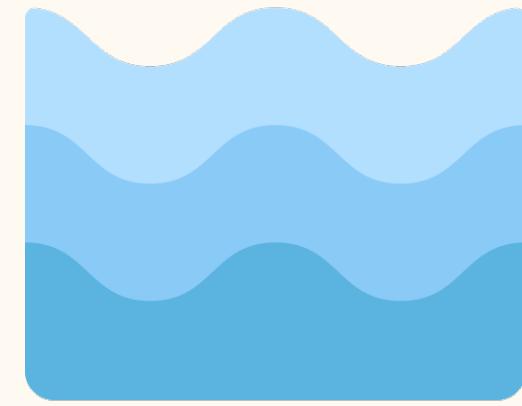
Atmospheric
composition



Waves



Sea ice



Ocean

- Current models are impressive, but **limited to one setting**.
- Unified approach?

Aurora



pretraining

- Train a single neural network a *large* body of Earth system data
- Learn general-purpose representation of dynamics that govern atmospheric and oceanic flow
- Slow and data hungry

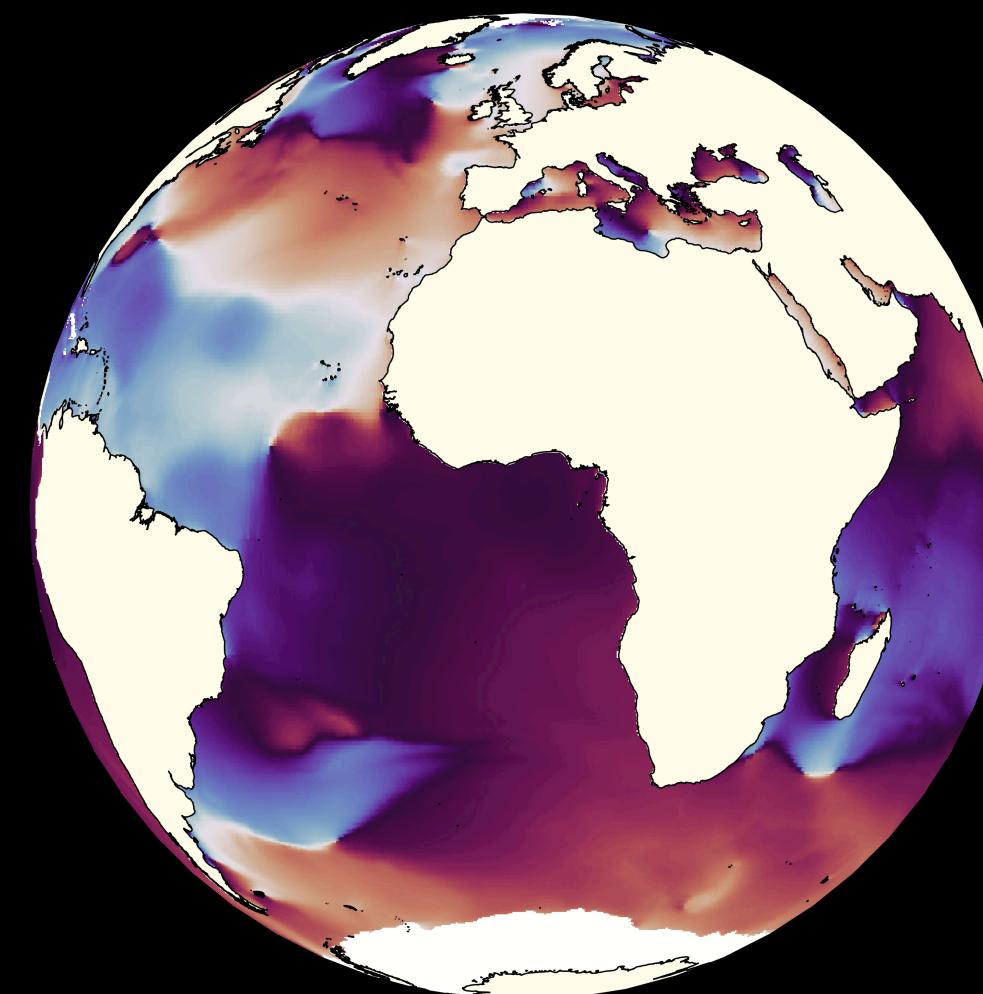
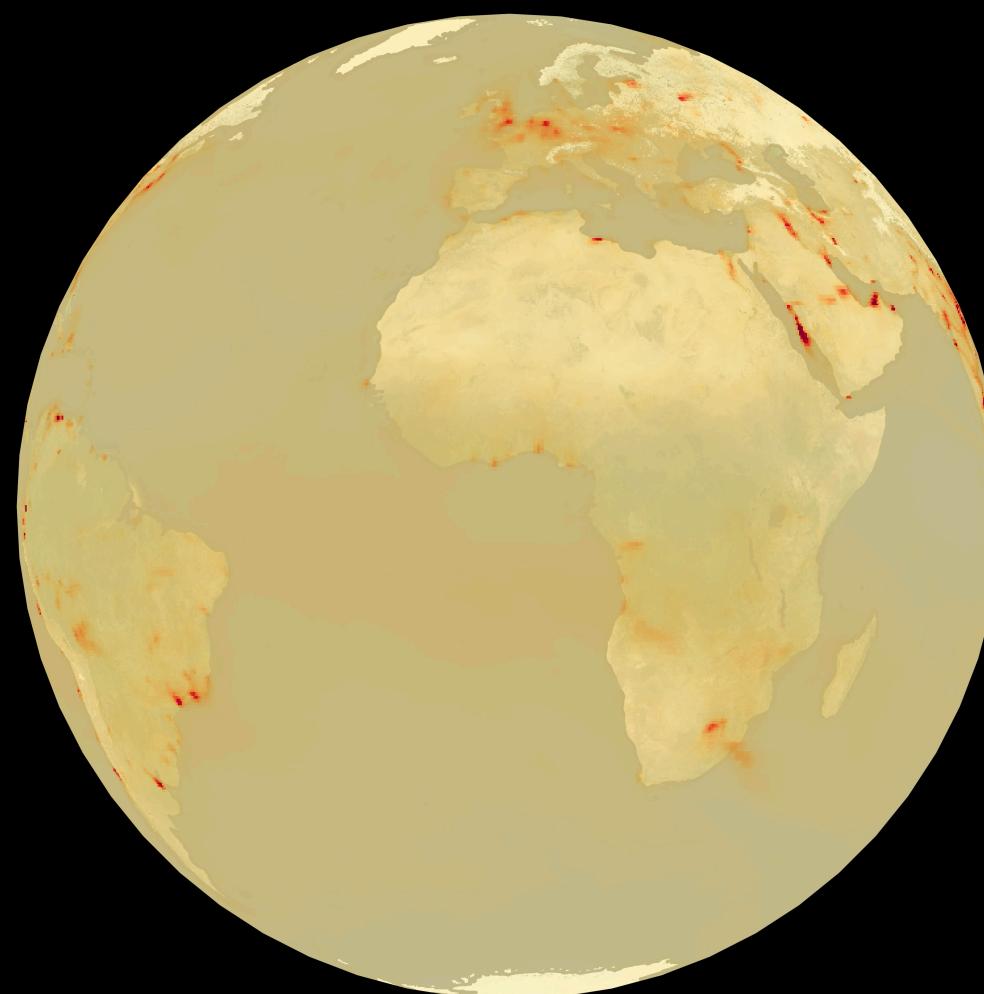
fine-tuning

- Leverage learned representation to **efficiently adapt to new domains!**
- Fast and data efficient

Aurora: a **foundation model** for the Earth system

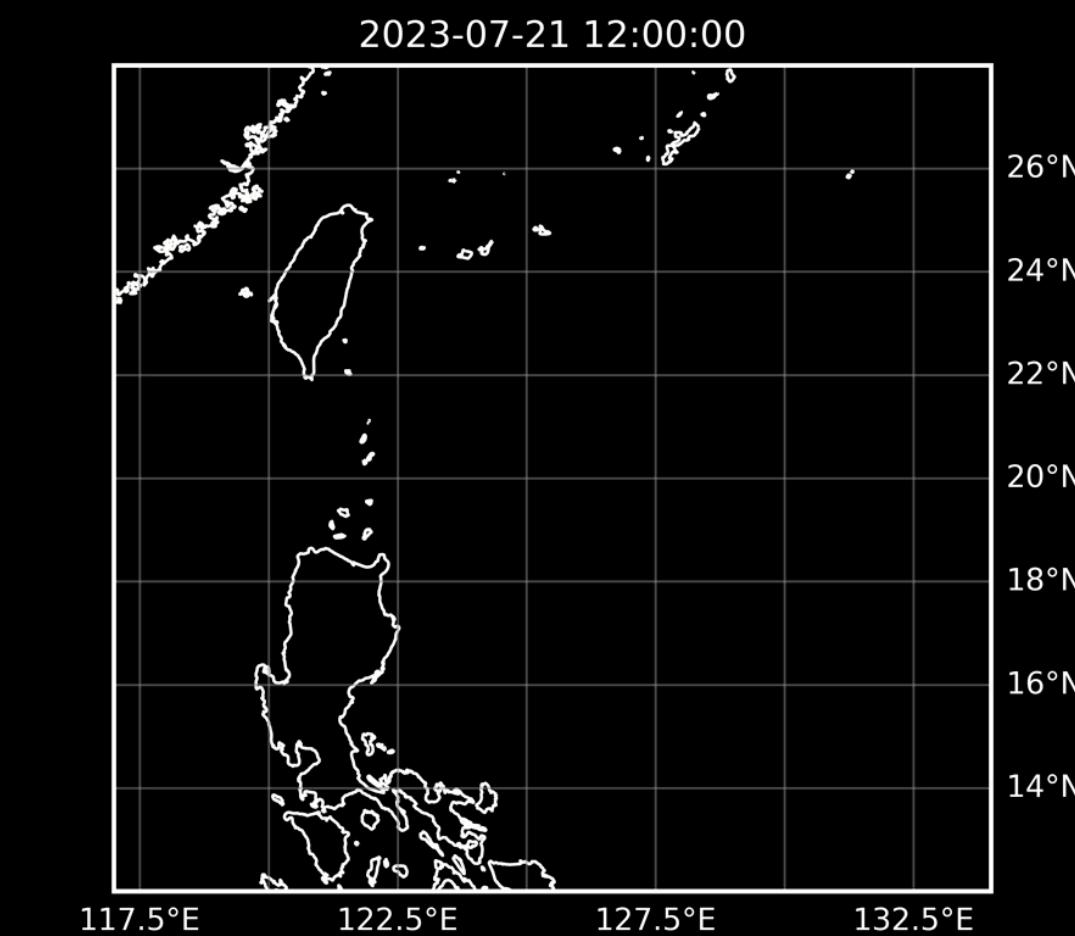
Aurora

Fine-Tuning Applications

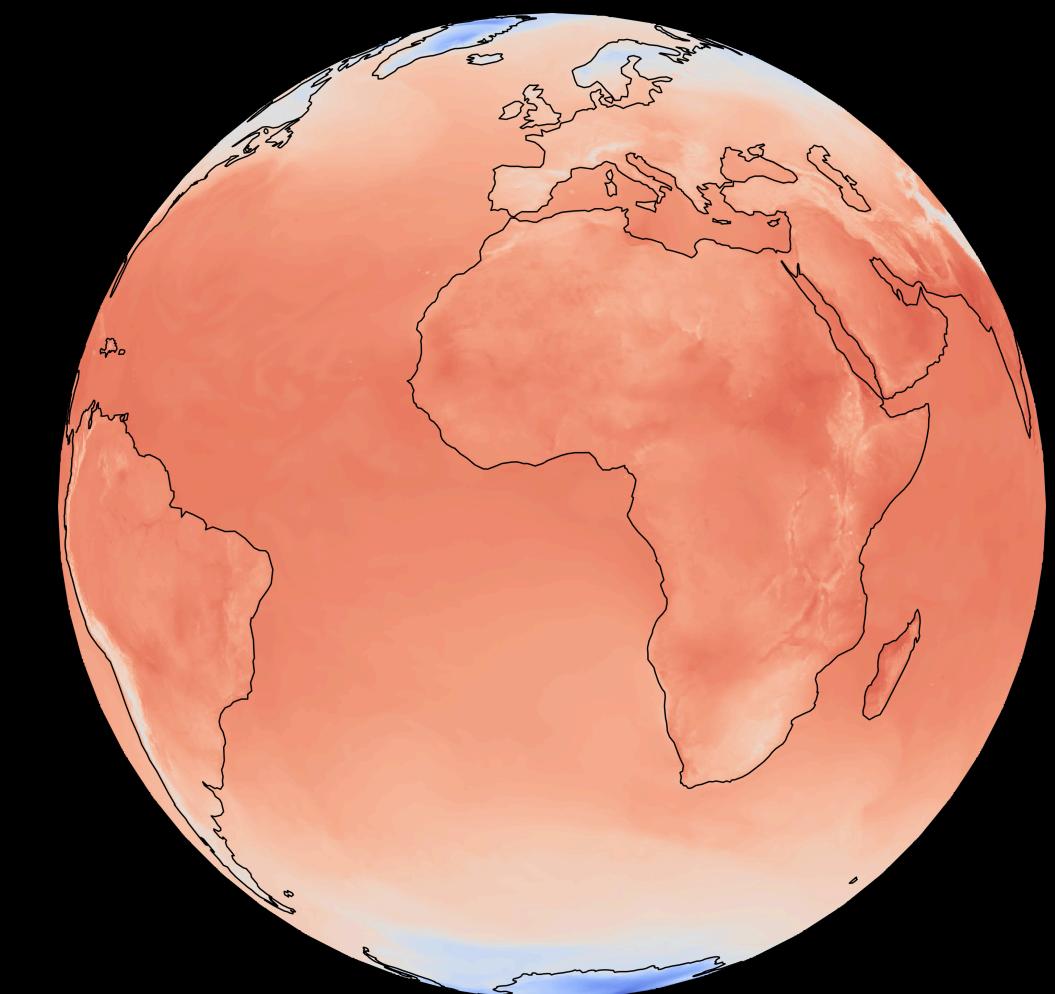


Atmospheric comp.
and air pollution

Ocean
waves



Tropical cyclone
track



High-resolution
weather

Aurora

Air Pollution Forecasting



- **Setup:** model levels of PM₁, PM_{2.5}, PM₁₀, CO, NO, NO₂, SO₂, O₃
- **Data:** Copernicus Atmospheric Monitoring Service (CAMS) analysis
- **Baseline:** CAMS forecasts

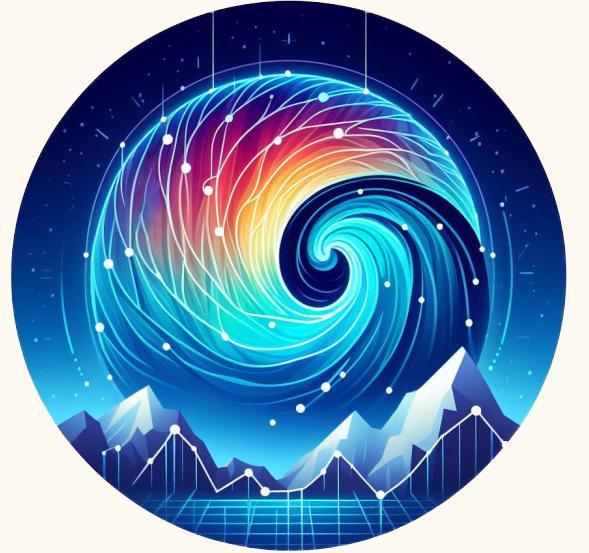


Coupled to IFS, ~10x more expensive:
~16 node-hours per hour lead time!

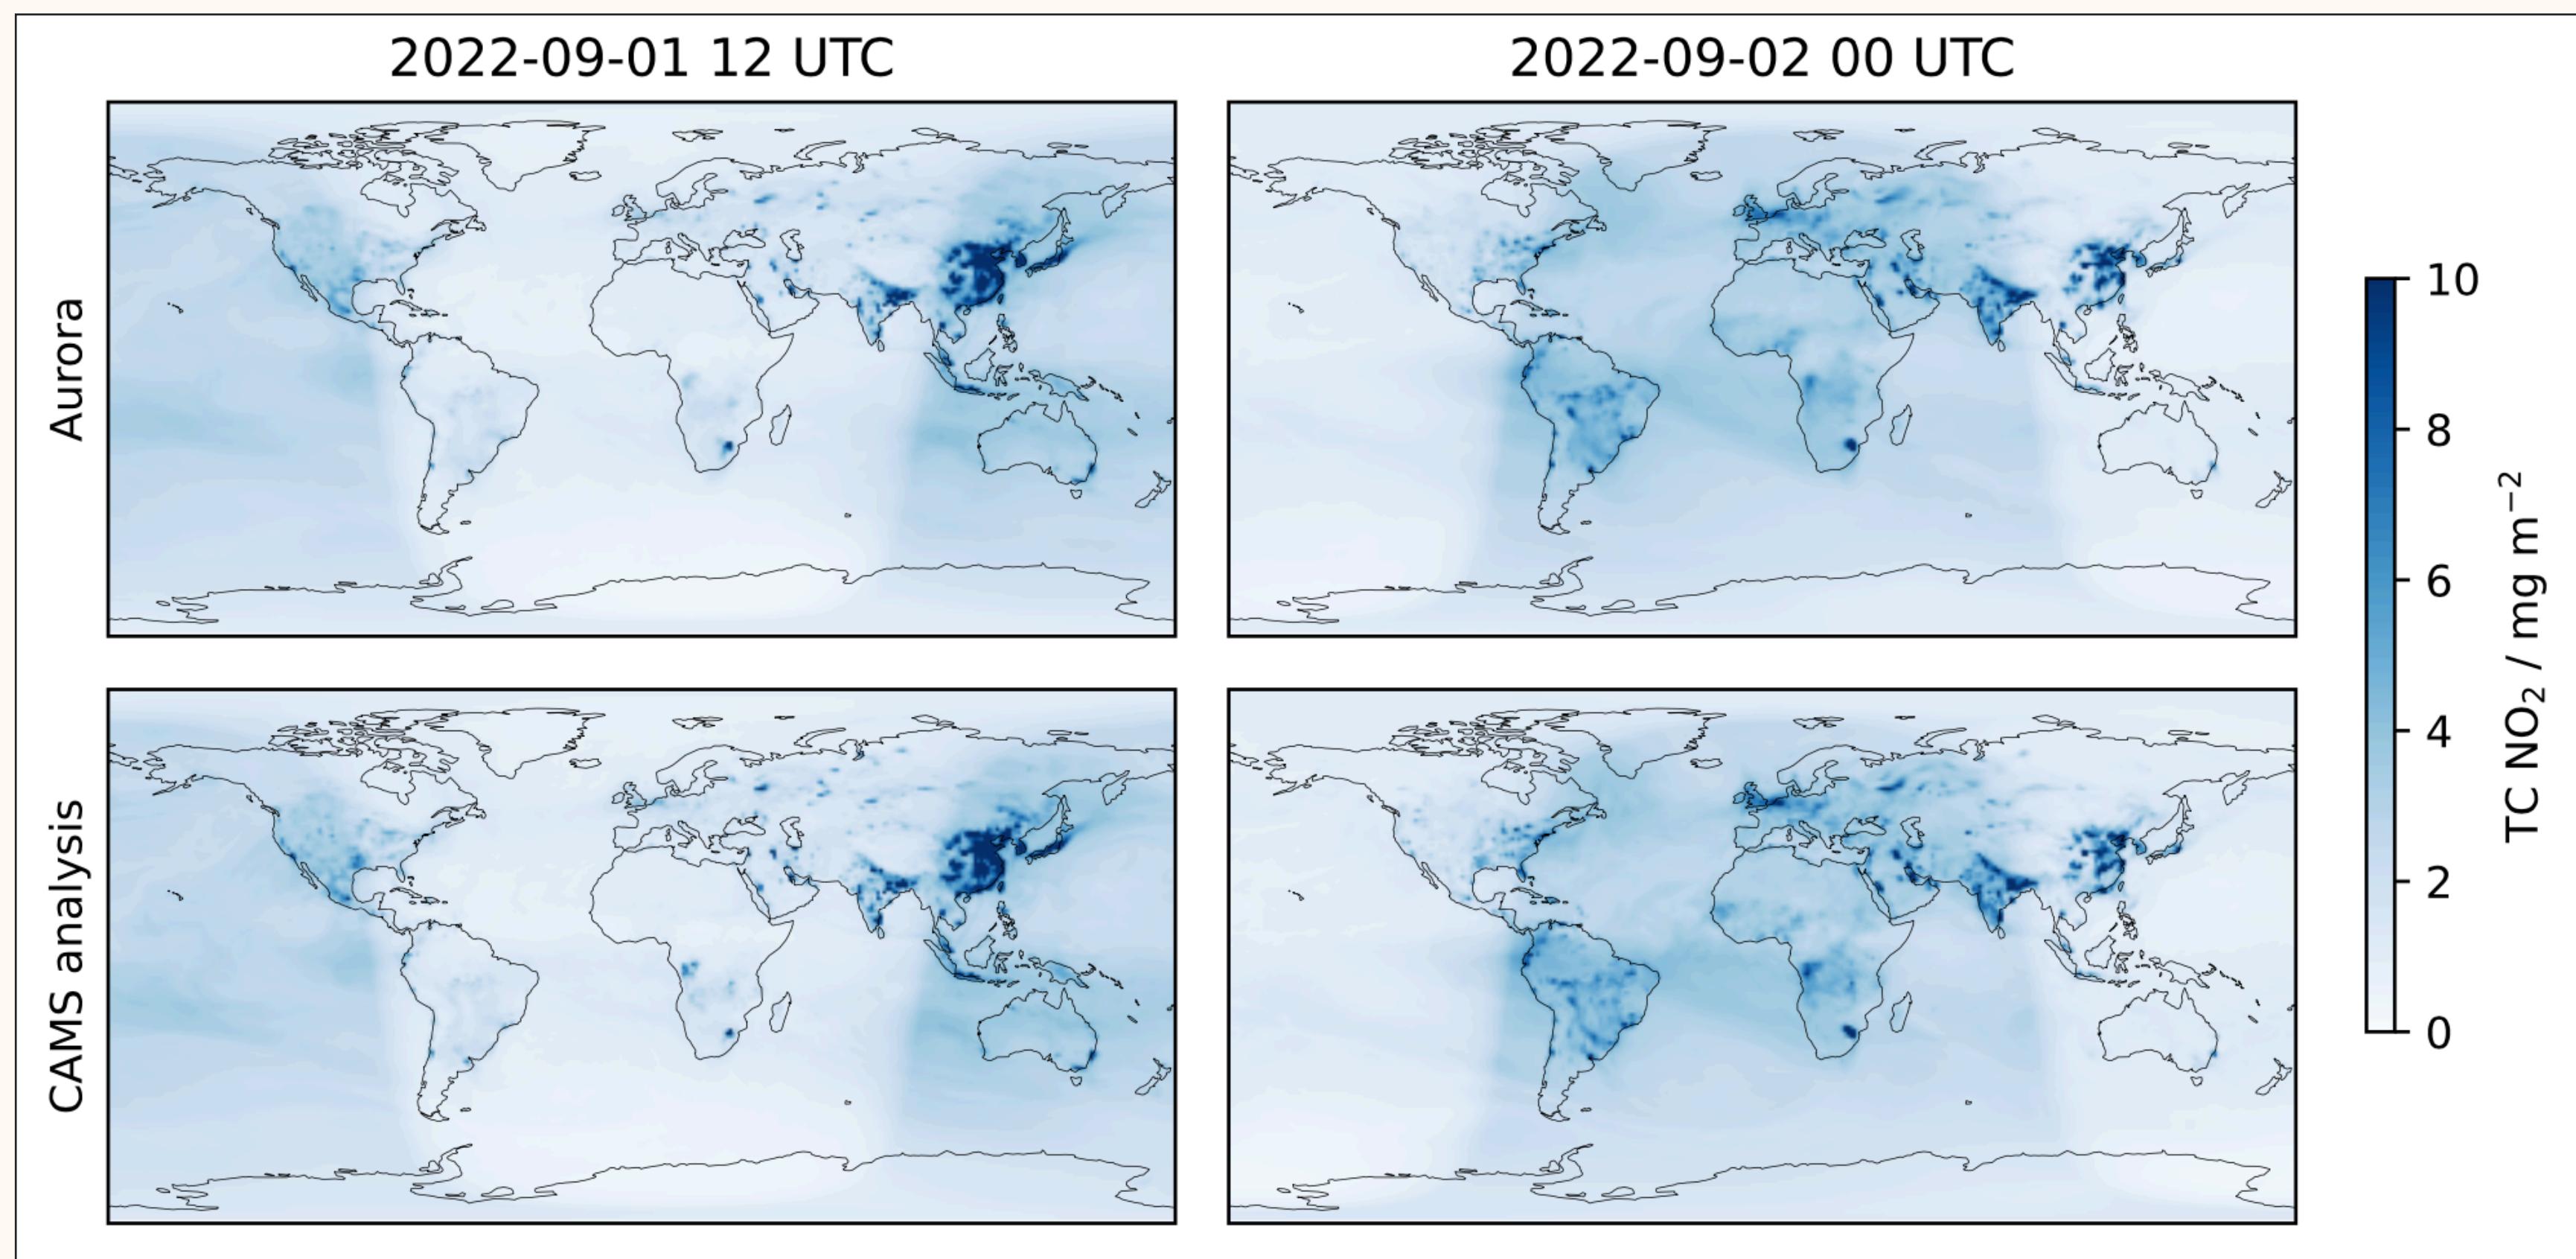
Aurora: **~0.5 s per hour lead time**

Aurora

Air Pollution Forecasting (2)



- Heterogeneous and spiky
- Anthropogenic factors
- Scarce
- Non-stationary



Overall:

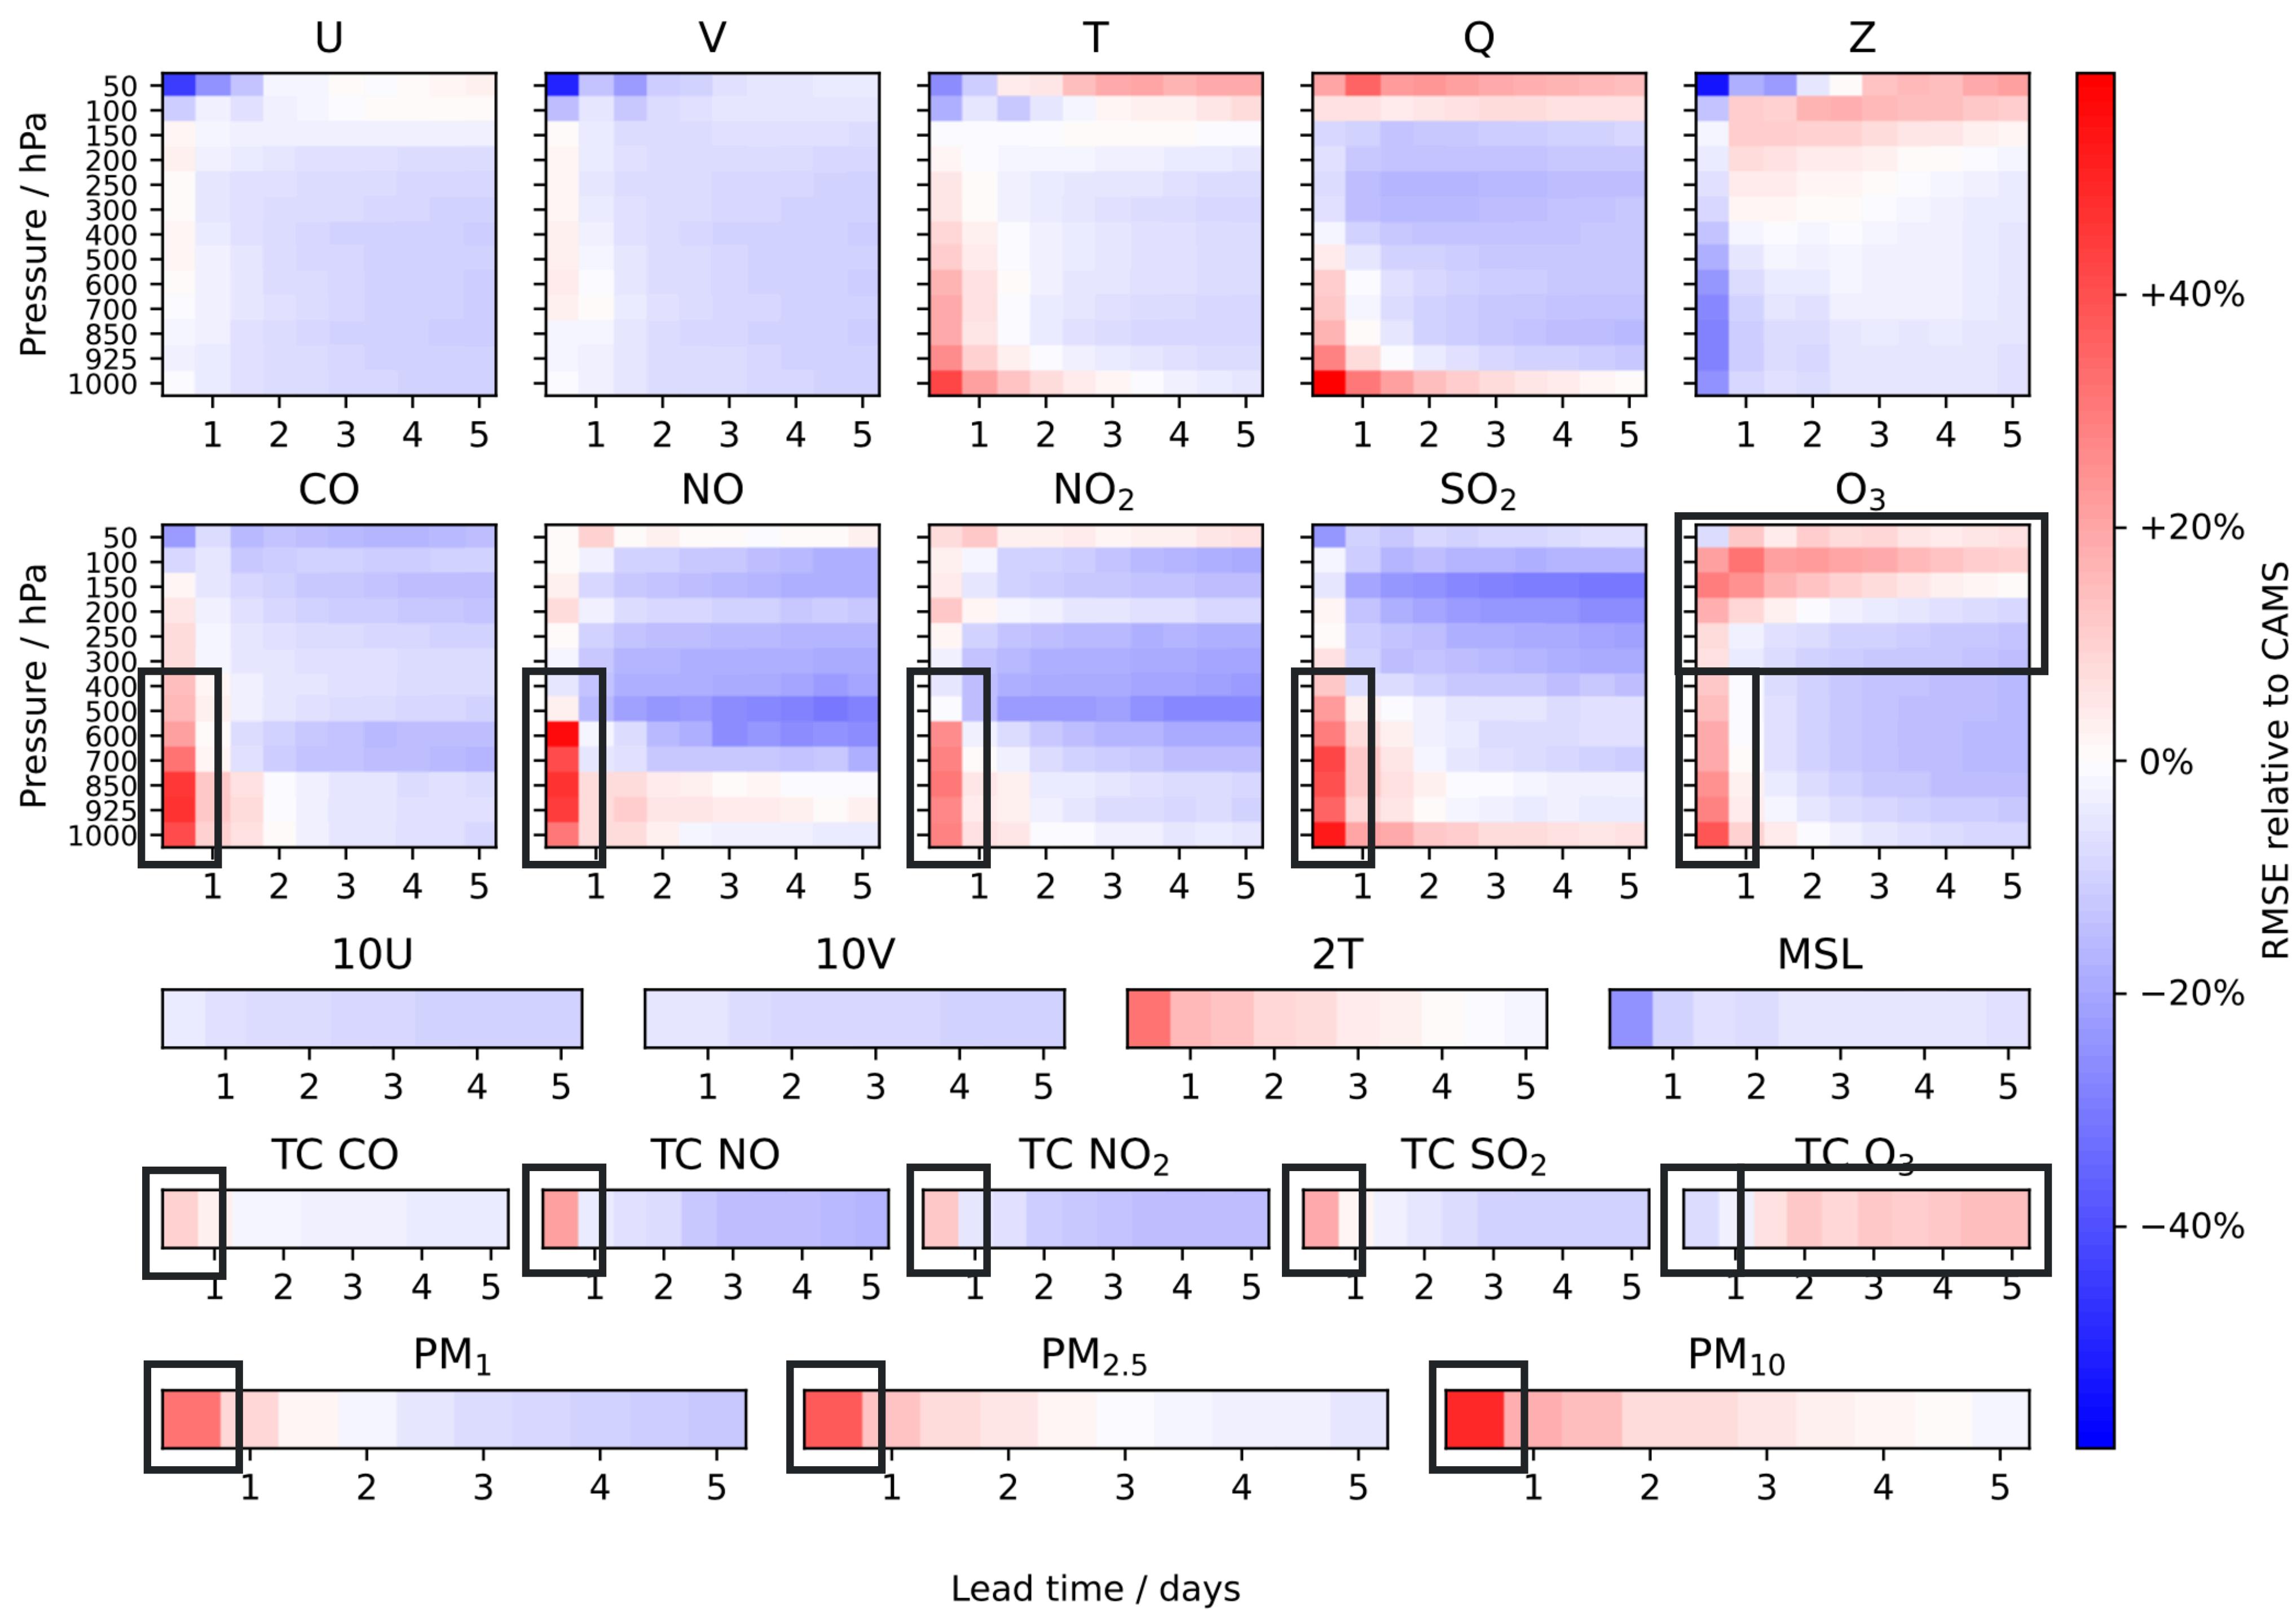
Competitive on
95%
(≤ 20% RMSE)

Better on 75%

Three days:

Competitive on
100%
(≤ 20% RMSE)

Better on 86%



Conclusion

- Medium-term weather forecasting has seen incredible progress
- **Pretraining–fine-tuning paradigm** to extend these advancements to other domains
- Aurora only scratches the surface!



wessel.ai/pdf/aurora



hi@wessel.ai