Spike and Slab Priors

1 Introduction

A spike and slab prior for a random variable X is a generative model—i.e., a prior—in which X either attains some fixed value v, called *the spike*, or is drawn some other prior $p_{\text{slab}}(x)$, called *the slab*. In the case that $v = 0$, X is either zero, or drawn from some other prior; in this case, the spike and slab prior is *sparsity inducing*, offering a principled alternative to e.g. sparsity-inducing regularisers.

The usual way of constructing a spike and slab prior is to introduce a latent variable $Z \sim \text{Ber}(\theta)$ where $Z = 0$ means that X attains the fixed value v and $Z = 1$ means that X is drawn from the slab $p_{\text{slab}}(x)$:

$$\begin{aligned} Z &\sim \text{Ber}(\theta), \\ X \mid Z = 0 &\sim \delta(x - v), \\ X \mid Z = 1 &\sim p_{\text{slab}}(x). \end{aligned}$$

Marginalising over Z, we equivalently have that

$$X \sim \theta p_X(x) + (1 - \theta)\delta(x - v),$$

which we recognise as a mixture model with mixture components $p_X(x)$ and $\delta(x - v)$, respectively having weights θ and $1 - \theta$. Figure 1 illustrates $p(x)$ in the case of a Gaussian slab.

2 Linear Regression with a Spike and Slab Prior

Let Y be an $\mathbb{R}$-valued random variable representing our observations at some point $x \in \mathbb{R}^n$, and consider the usual model for linear regression:

$$Y \mid \beta, x \sim \mathcal{N}(\langle \beta, x \rangle, \sigma^2)$$

where $\langle \beta, x \rangle$ denotes the inner product between β and x. In the case that x is high dimensional, we might not have enough data to accurately estimate the coefficients β. One way to mitigate this issue is to build zeros into β, and putting a spike and slab prior on β is a perfectly viable

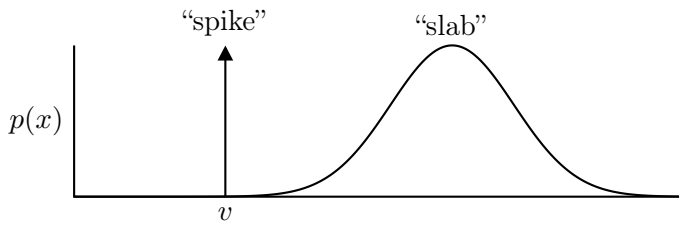

Figure 1: The density $p(x)$ in the case of a Gaussian slab

approach to do so:

$$\begin{aligned} Z_i &\sim \mathrm{Ber}(\theta), \\ \beta_i \,|\, Z_i = 0 &\sim \delta(\beta_i), \\ \beta_i \,|\, Z_i = 1 &\sim \mathcal{N}(0, \tau^{-1}). \end{aligned}$$

We can equivalently formulate the resulting model in a slightly more compact and convenient form:

$$\begin{aligned} Z_i &\sim \mathrm{Ber}(\theta), \\ \beta_i &\sim \mathcal{N}(0, \tau^{-1}), \\ Y \,|\, Z, \beta, x &\sim \mathcal{N}(\langle z \circ \beta, x \rangle, \sigma^2) \end{aligned}$$

where $\circ$ denotes the Hadamard product. Indeed, $(z \circ \beta)_i = z_i \beta_i = 0$ if $z_i = 0$ and similarly $(z \circ \beta)_i = \beta_i$ if $z_i = 1$.

Upon observing data $\mathcal{D} = (x^{(t)}, y^{(t)})_{t=1}^T$, we wish to compute our posterior belief about β and Z:

$$p(z, \beta \,|\, \mathcal{D}) = \frac{1}{p(\mathcal{D})} p(z) p(\beta) \prod_{t=1}^{T} p(y^{(t)} \,|\, z, \beta, x^{(t)})$$

where $p(\mathcal{D})$ denotes the *evidence*:

$$p(\mathcal{D}) = \int p(z) p(\beta) \prod_{t=1}^{T} p(y^{(t)} \,|\, z, \beta, x^{(t)}) \,\mathrm{d}z \,\mathrm{d}\beta.$$

Unfortunately, $p(\mathcal{D})$ is hard to compute, because it requires summing over the 2^n values that Z can attain, and it is not clear how to efficiently do so. We must therefore resort to *approximate inference*.

To perform inference in models employing a spike and slab prior, a sampling-based approach, *Gibbs sampling* in particular, is often used. Gibbs sampling states that given some initial $Z^{(0)}$

and $\beta^{(0)}$, iterating

$$\begin{aligned}
Z_1^{(i)} &\sim p(z_1 \,|\, z_{2:n}^{(i-1)}, \beta^{(i-1)}, \mathcal{D}),\\
Z_2^{(i)} &\sim p(z_2 \,|\, z_1^{(i)}, z_{3:n}^{(i-1)}, \beta^{(i-1)}, \mathcal{D}),\\
&\vdots\\
Z_n^{(i)} &\sim p(z_n \,|\, z_{1:n-1}^{(i)}, \beta^{(i-1)}, \mathcal{D}),\\
\beta^{(i)} &\sim p(\beta \,|\, z^{(i)}, \mathcal{D})
\end{aligned}$$

will eventually yield samples from the joint posterior:

$$(Z^{(i)}, \beta^{(i)}) \sim p(z, \beta \,|\, \mathcal{D}) \quad \text{for large enough } i.$$

Fortunately, these conditionals are easy to compute[1]:[✓]

$$\begin{aligned}
\log p(z_i \,|\, z_{-i}, \beta, \mathcal{D}) &\simeq \log p(z_i) + \log p(\mathcal{D} \,|\, z, \beta)\\
&\simeq \log p(z_i) + \frac{1}{\sigma^2}\langle z \circ \beta, \hat{\mu}\rangle - \frac{1}{2\sigma^2}\langle z \circ \beta, \hat{\Sigma}(z \circ \beta)\rangle,\\
\hat{\Sigma} &= \sum_{t=1}^{T} x^{(t)} x^{(t)\mathsf{T}},\\
\hat{\mu} &= \sum_{t=1}^{T} x^{(t)} y^{(t)},
\end{aligned}$$

and

$$\begin{aligned}
p(\beta \,|\, z, \mathcal{D}) &\propto p(\beta) p(\mathcal{D} \,|\, z, \beta)\\
&\propto \mathcal{N}(\beta; (\tau\sigma^2 I + \tilde{\Sigma})^{-1}\tilde{\mu}, \sigma^2(\tau\sigma^2 I + \tilde{\Sigma})^{-1})\\
\tilde{\Sigma} &= \sum_{t=1}^{T} (z \circ x^{(t)})(z \circ x^{(t)})^{\mathsf{T}},\\
\tilde{\mu} &= \sum_{t=1}^{T} z \circ x^{(t)} y^{(t)}.
\end{aligned}$$

One can now sample and happily compute expectations under the posterior distribution.

Remark 2.1. Although the generative model specifies each weight β_i to be either zero or nonzero, the posterior over Z_i will not conclude either case: the posterior over Z_i instead assigns probabilities to both possibilities of β_i being zero or nonzero. Therefore, the posterior

[1] In the case that the conditionals cannot be computed analytically, one could use another MCMC method, often Metropolis–Hastings, to sample from the conditionals, yielding a composite procedure often referred to as *Metropolis–Hastings within Gibbs*.

distribution does not yield a "sparse solution", but rather a weighting of all possible sparse solutions. And this makes perfect sense: only in the limit of infinite data can the model conclude a weight to be zero. ◁

3 Diagnosing MCMC

3.1 Big O Notation and Convergence of Simple Monte Carlo Estimates

To begin with, let us quickly recap how the number of samples relates to the accuracy of a Monte Carlo estimate.

Definition 3.1 (Small O: Convergence in Probability)**.** If (X_n) is a sequence of random variables and (a_n) a sequence of constants, then $X_n = o_P(a_n)$ means that for every $\varepsilon > 0$,

$$\lim_{n\to\infty} P(|X_n/a_n| < \varepsilon) = 1.$$

◁

If $X_n = o_P(a_n)$, then that X_n will eventually become arbitrarily close to a_n in probability; in other words, asymptotically X_n behaves like a_n.

Definition 3.2 (Big O: Stochastic Boundedness)**.** If (X_n) is a sequence of random variables and (a_n) a sequence of constants, then $X_n = O_P(a_n)$ means that for every $\varepsilon > 0$ there exist $\delta_\varepsilon > 0$ and $N_\varepsilon > 0$ such that

$$P(|X_n/a_n| \leq \delta_\varepsilon) \geq 1 - \varepsilon \quad \text{for all } n \geq N_\varepsilon.$$

◁

If $X_n = O_P(a_n)$, then that means that for every $\varepsilon > 0$, we can identify a region around a_n that eventually will contain X_n with probability at least $1 - \varepsilon$; in other words, asymptotically X_n is finitely far from a_n.

Now, consider the simple Monte Carlo estimator

$$X_n = \frac{1}{n}\sum_{i=1}^{n} X_i,$$

where all X_i are drawn i.i.d. Let $\mathbb{E}X_i = \mu$ and regard $\mathbb{V}X_i = \sigma^2$ as a constant. We then find that

Proposition 3.1. If (X_i) are drawn i.i.d., then $\frac{1}{n}\sum_{i=1}^{n} X_i - \mu = O_p(1/\sqrt{n})$. ◁

In other words, to gain a digit more accurate results, one needs 100 times more samples.

3.2 Convergence of MCMC Estimates

Unfortunately, the analysis in the foregoing section does not apply to averages computed with samples from a Markov chain: in that case, the samples are not i.i.d., but *correlated* instead.

Definition 3.3 (Effective Sample Size (ESS))**.** For n samples with autocorrelation ρ_t, the *effective sample size* is $n_{\text{ESS}} = n/\tau$ where

$$\tau = 1 + 2\sum_{i=1}^{\infty} \rho_i.$$

◁

For the purpose of computing averages, the effective sample size n_{ESS} is the number independent samples "contained" in correlated samples from a Markov chain. Using this number of "effective samples" n_{ESS}, one can apply the analysis from the foregoing section.

Proposition 3.2. If (X_i) are drawn from a Markov chain, then

$$\frac{1}{n}\sum_{i=1}^{n} X_i - \mu = O_p(1/\sqrt{n_{\text{ESS}}}).$$

◁

In other words, to gain a digit more accurate results, one needs 100 times more effective samples.

3.3 Geweke Test

Testing MCMC code is often difficult (Grosse and Duvenaud, 2014): algorithms are stochastic, algorithms may perform badly for reasons other than incorrect implementations, and good performance is often a matter of judgement.

We discuss one way of testing correctness of an MCMC algorithm, a test called the *Geweke test* (Geweke, 2004). Sampling from a joint distribution $p(x, z)$ can be done in two different ways:

(1) Sample (X, Z) from the generative model: first sample $Z \sim p(z)$ and then $X \mid Z \sim p(x \mid z)$.

(2) Start with a sample $X \sim p(x)$ from the generative model, and produce a sample from $Z \mid X \sim p(z \mid x)$ using the MCMC algorithm. Finally, resample $X \mid Z \sim p(x \mid x)$.

The Geweke test tests that (1) and (2) produce samples from the same distribution. To do this, one can follow Grosse and Duvenaud (2014) and simply employ a P–P plot; Figure 2 illustrates a negative result, an indeterminate result, and a positive result.

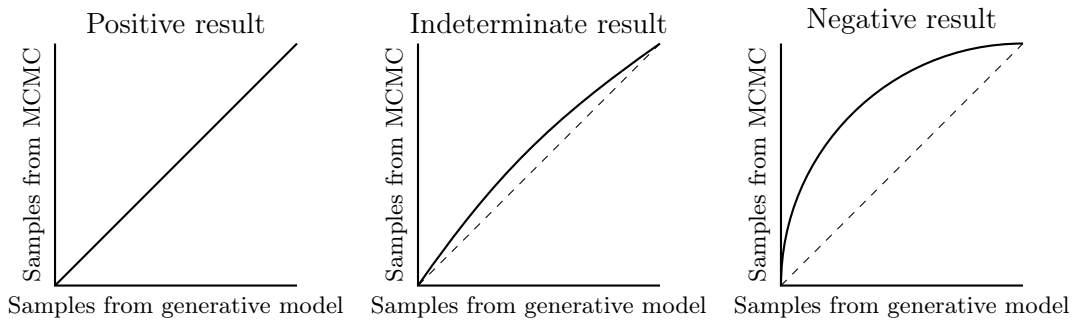

Figure 2: Various outcomes of the P–P in a Geweke test

References

Geweke, J. (2004). Getting it right: Joint distribution tests of posterior simulators. *Journal of the American Statistical Association*, *99*(467), 799–804. (Cit. on p. 5).

Grosse, R. B., & Duvenaud, D. K. (2014). Testing MCMC code. *arXiv preprint arXiv:1412.5218*. eprint: https://arxiv.org/abs/1412.5218. (Cit. on pp. 5, 6)